Lecture 11

Weighted Graphs: Dijkstra and Bellman-Ford

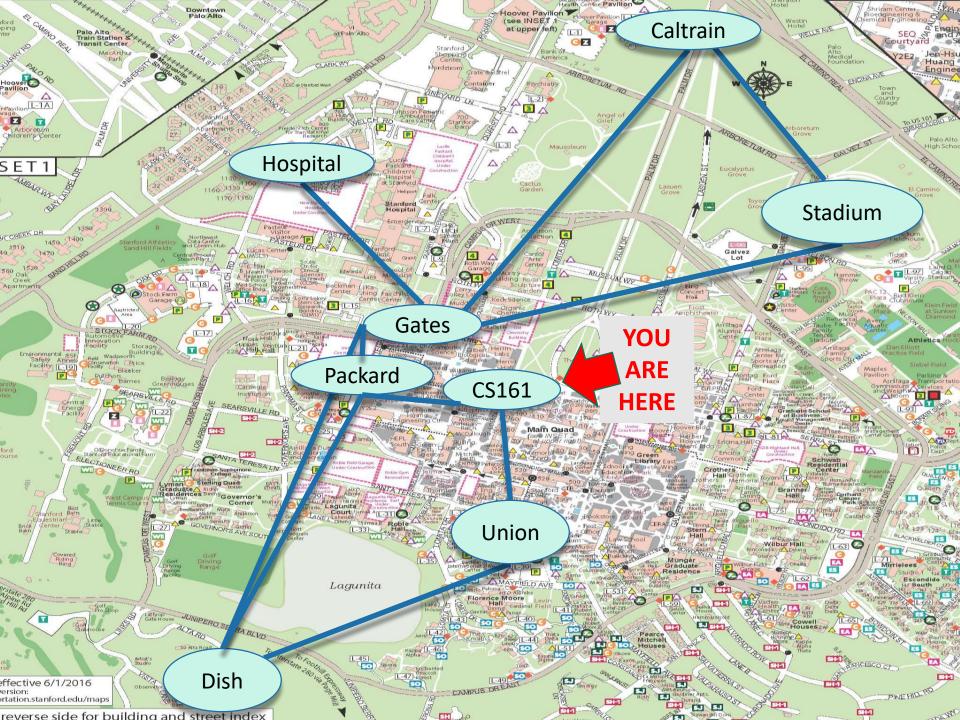
Last Lecture

- Graphs!
- DFS
 - Topological Sorting
 - Strongly Connected Components
- BFS
 - Shortest Paths in unweighted graphs

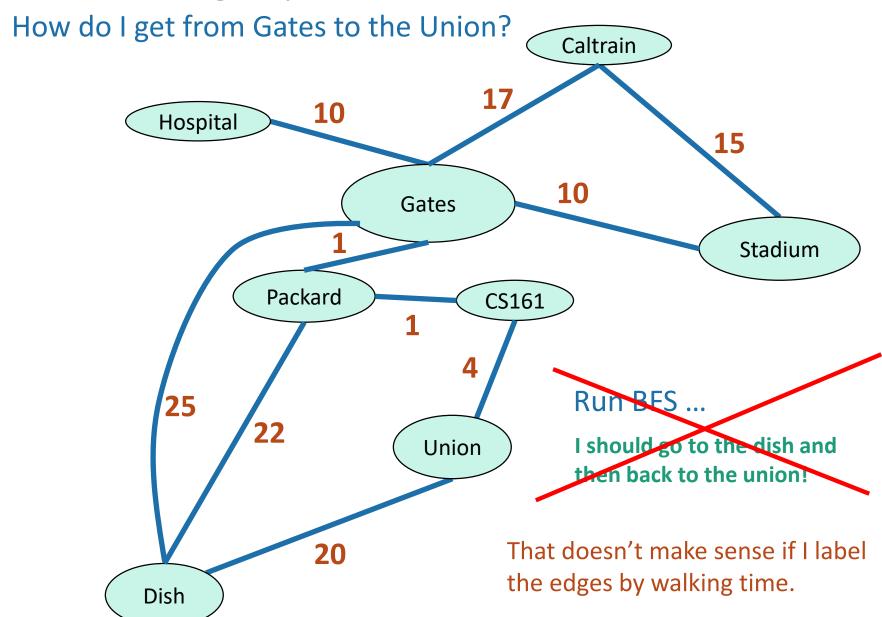
Today

- What if the graphs are weighted?
 - All nonnegative weights: Dijkstra!
 - If there are negative weights: Bellman-Ford!

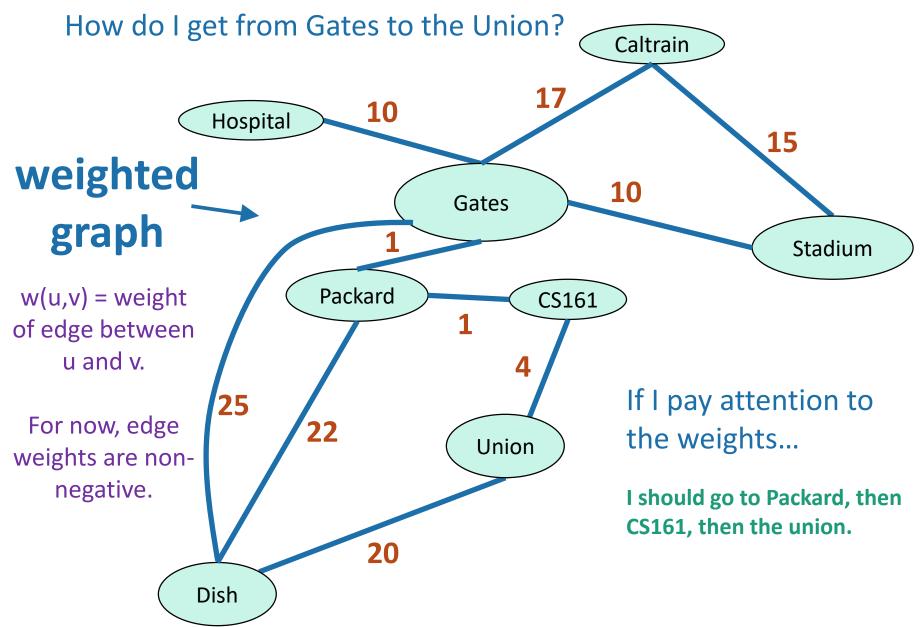




Just the graph

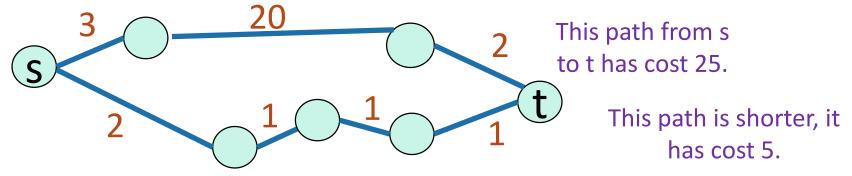


Just the graph

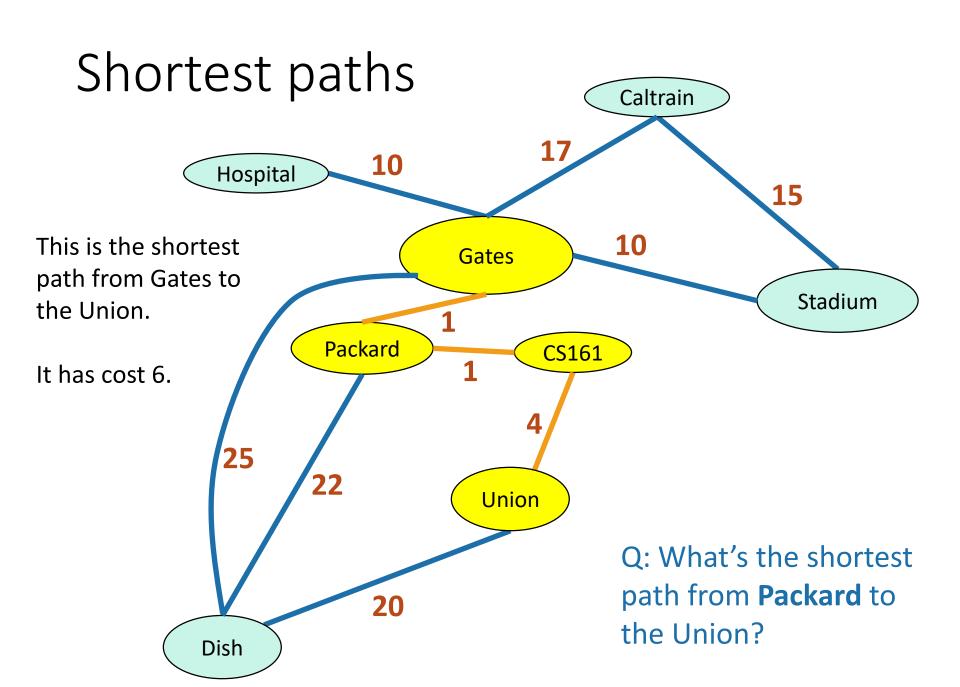


Shortest path problem

- What is the shortest path between u and v in a weighted graph?
 - the cost of a path is the sum of the weights along that path
 - The shortest path is the one with the minimum cost.



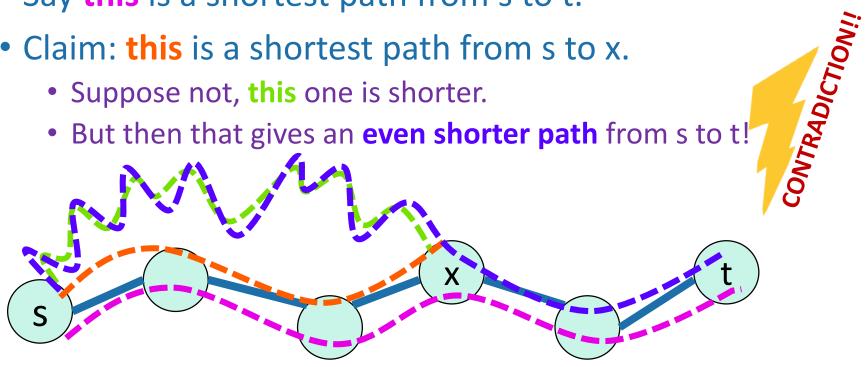
- The distance d(u,v) between two vertices u and v is the cost of the the shortest path between u and v.
- For this lecture all graphs are directed, but to save on notation I'm just going to draw undirected edges.



Warm-up

A sub-path of a shortest path is also a shortest path.

- Say this is a shortest path from s to t.
- Claim: this is a shortest path from s to x.
 - Suppose not, this one is shorter.
 - But then that gives an even shorter path from s to t!



Single-source shortest-path problem

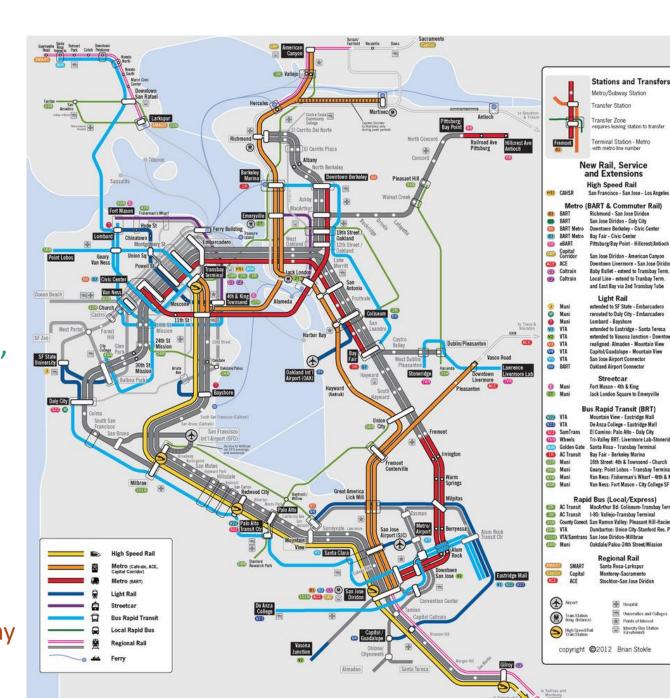
• I want to know the shortest path from one vertex (Gates) to all other vertices.

Destination	Cost	To get there
Packard	1	Packard
CS161	2	Packard-CS161
Hospital	10	Hospital
Caltrain	17	Caltrain
Union	6	Packard-CS161-Union
Stadium	10	Stadium
Dish	23	Packard-Dish

(Not necessarily stored as a table – how this information is represented will depend on the application)

Example

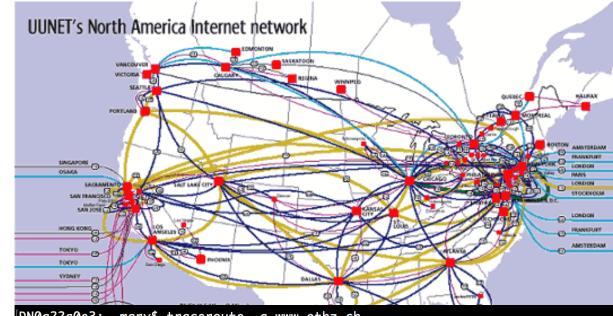
- I regularly have to solve "what is the shortest path from Palo Alto to [anywhere else]" using BART, Caltrain, lightrail, MUNI, bus, Amtrak, bike, walking, uber/lyft.
- Edge weights have something to do with time, money, hassle. (They also change depending on my mood and traffic...).



Example

Network routing

- I send information over the internet, from my computer to to all over the world.
- Each path (from a router to another router) has a cost which depends on link length, traffic, other costs, etc..
- How should we send packets?



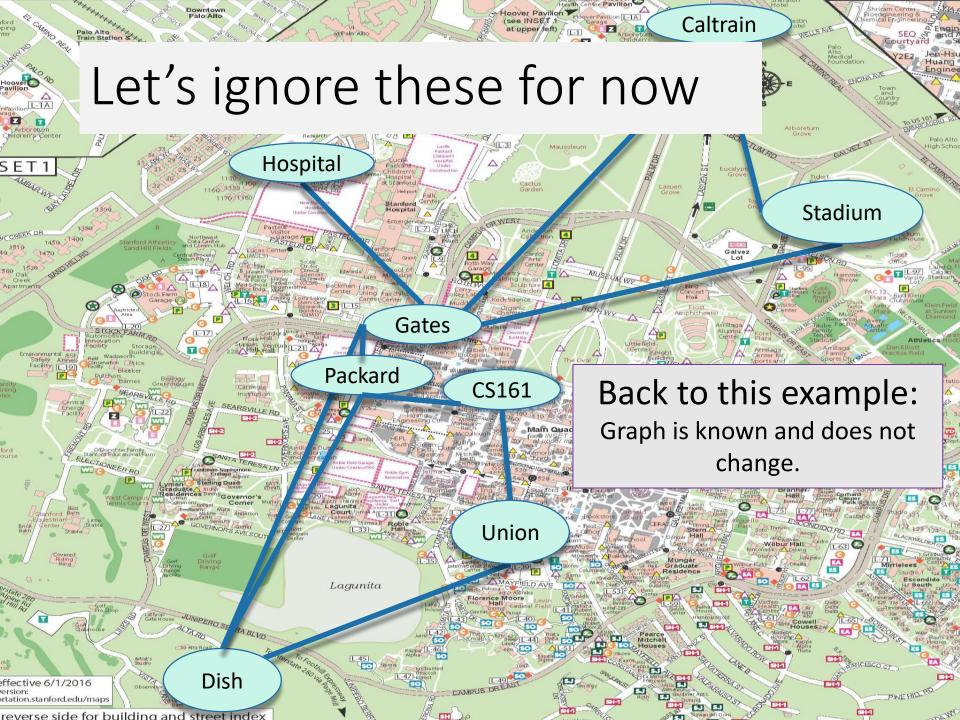
```
DN0a22a0e3:~ mary$ traceroute -a www.ethz.ch
traceroute to www.ethz.ch (129.132.19.216), 64 hops max, 52 byte packets
   [ASO] 10.34.160.2 (10.34.160.2) 38.168 ms 31.272 ms 28.841 ms
   [AS0] cwa-vrtr.sunet (10.21.196.28) 33.769 ms 28.245 ms 24.373 ms
   [AS32] 171.66.2.229 (171.66.2.229) 24.468 ms 20.115 ms 23.223 ms
   [AS32] hpr-svl-rtr-vlan8.sunet (171.64.255.235) 24.644 ms 24.962 ms 1
   [AS2152] hpr-svl-hpr2--stan-ge.cenic.net (137.164.27.161) 22.129 ms 4.9
   [AS2152] hpr-lax-hpr3--svl-hpr3-100ge.cenic.net (137.164.25.73) 12.125 |
   [AS2152] hpr-i2--lax-hpr2-r&e.cenic.net (137.164.26.201) 40.174 ms 38.3
   [ASO] et-4-0-0.4079.sdn-sw.lasv.net.internet2.edu (162.252.70.28) 46.57
   [AS0] et-5-1-0.4079.rtsw.salt.net.internet2.edu (162.252.70.31)
   [AS0] et-4-0-0.4079.sdn-sw.denv.net.internet2.edu (162.252.70.8)
                                                                   47.454
   [AS0] et-4-1-0.4079.rtsw.kans.net.internet2.edu (162.252.70.11)
                                                                   70.825 1
   [AS0] et-4-1-0.4070.rtsw.chic.net.internet2.edu (198.71.47.206)
                                                                   77.937 ı
   [ASO] et-0-1-0.4079.sdn-sw.ashb.net.internet2.edu (162.252.70.60) 77.68
   [AS0] et-4-1-0.4079.rtsw.wash.net.internet2.edu (162.252.70.65) 71.565 i
   [AS21320] internet2-qw.mx1.lon.uk.geant.net (62.40.124.44) 154.926 ms
   [AS21320] ae0.mx1.lon2.uk.geant.net (62.40.98.79) 146.565 ms 146.604 ms
   [AS21320] ae0.mx1.par.fr.geant.net (62.40.98.77) 153.289 ms 184.995 ms
   [AS21320] ae2.mx1.gen.ch.geant.net (62.40.98.153) 160.283 ms 160.104 ms
   [AS21320] swice1-100ge-0-3-0-1.switch.ch (62.40.124.22) 162.068 ms 160
   [AS559] swizh1-100ge-0-1-0-1.switch.ch (130.59.36.94) 165.824 ms 164.2
   [AS559] swiez3-100ge-0-1-0-4.switch.ch (130.59.38.109) 164.269 ms 164.3
   [AS559] rou-gw-lee-tengig-to-switch.ethz.ch (192.33.92.1) 164.082 ms 1
```

[AS559] rou-fw-rz-rz-gw.ethz.ch (192.33.92.169) 164.773 ms 165.193 ms

A few things that make these examples even more difficult

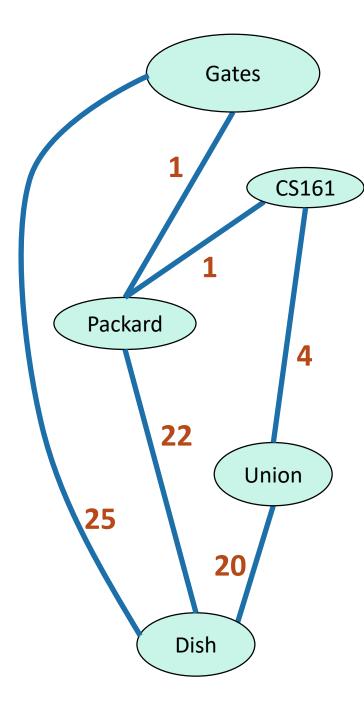
than trying to navigate the Stanford campus.

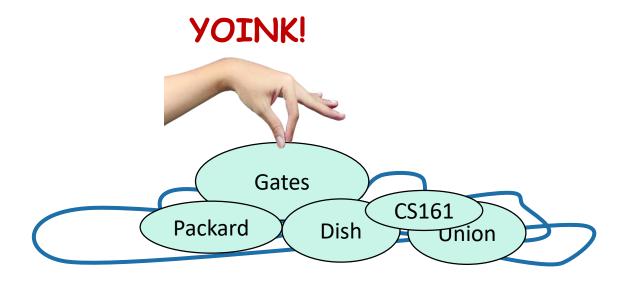
- Costs may change
 - If it's raining the cost of biking is higher
 - If a link is congested, the cost of routing a packet along it is higher
- The network might not be known
 - My computer doesn't store a map of the internet
- We want to do these tasks really quickly
 - More seriously, the internet.



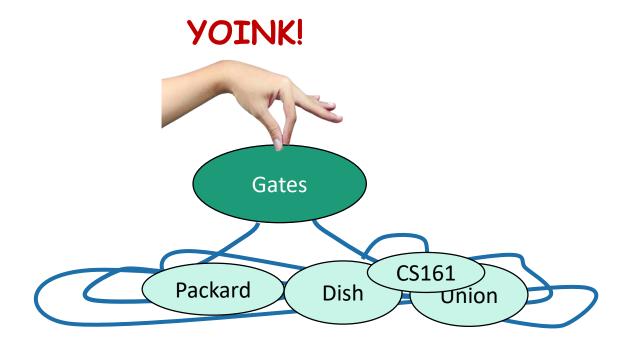
Dijkstra's algorithm

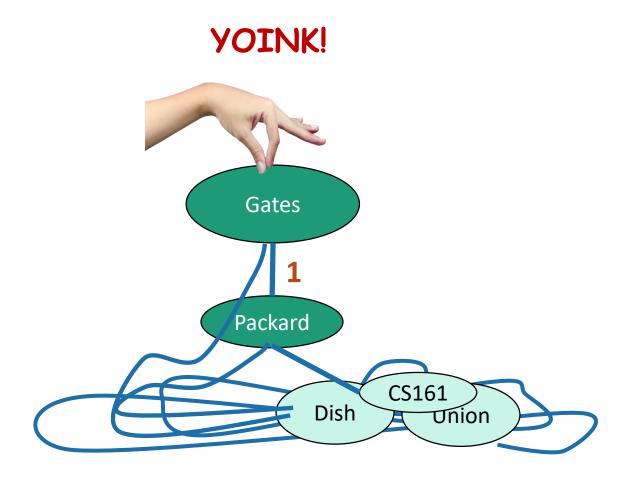
 What are the shortest paths from Gates to everywhere else?



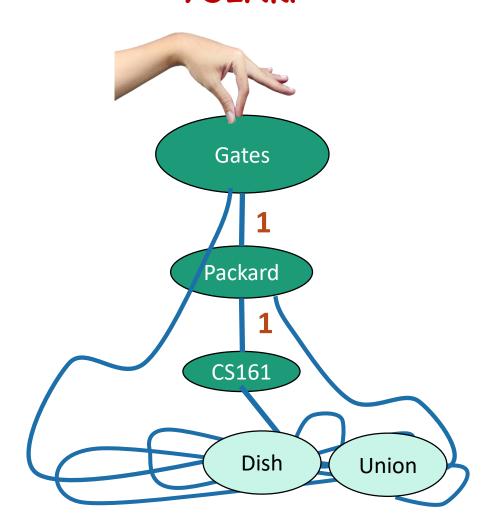



A vertex is done when it's not on the ground anymore.






YOINK!









This also creates a tree structure!

The shortest paths are the lengths along this tree.



How do we actually implement this?


Without string and gravity?









How far is a node from Gates?

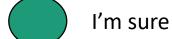


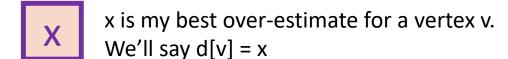
vertices v, and v[Gates] = 0

 Pick the not-sure node u with the smallest estimate d[u].






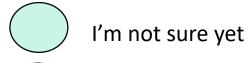


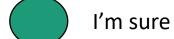


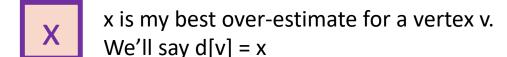


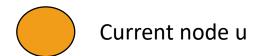
- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v], d[u] + edgeWeight(u,v))



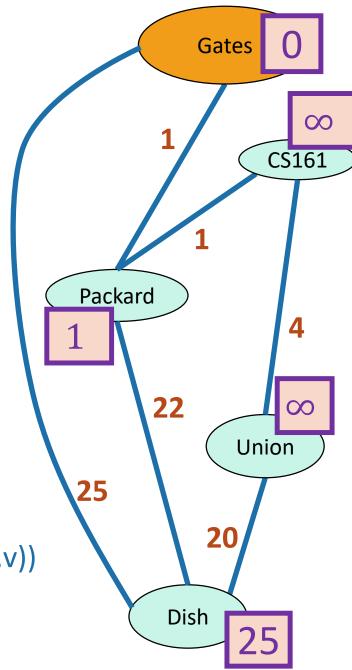




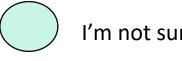




- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.



How far is a node from Gates?



I'm not sure yet



I'm sure

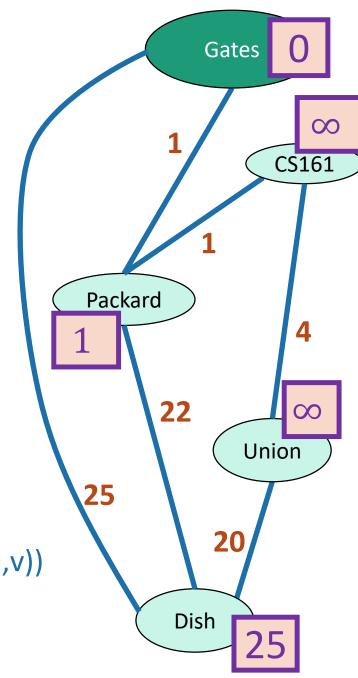


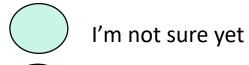
x is my best over-estimate for a vertex v. We'll say d[v] = x

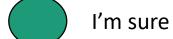


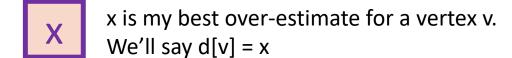
Current node u

- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as **SUCE**.
- Repeat



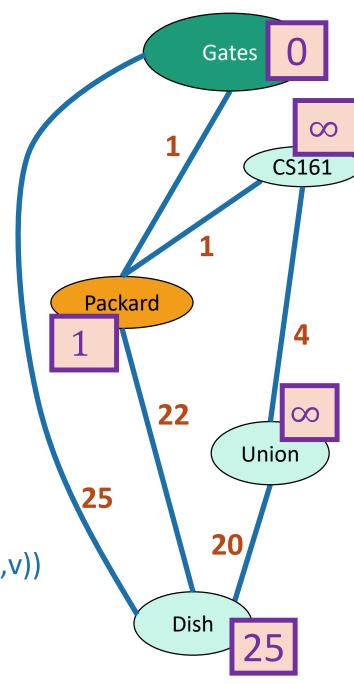


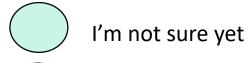


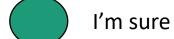


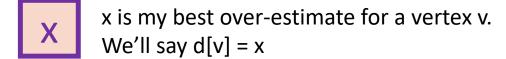


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



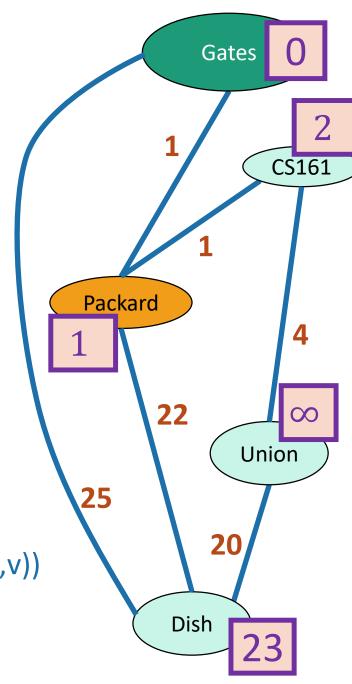


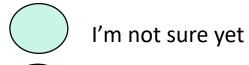


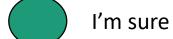


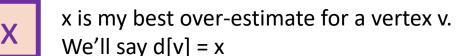


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



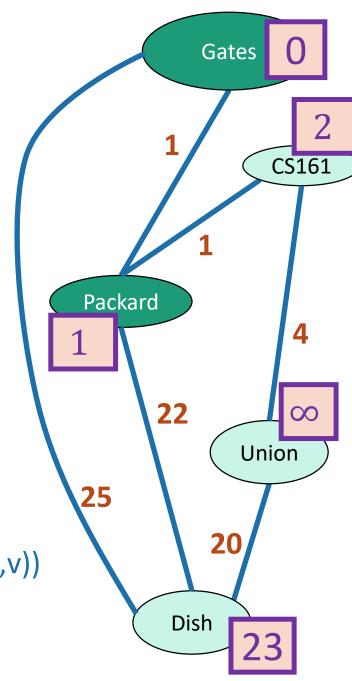


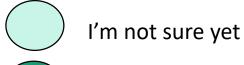


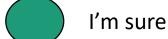


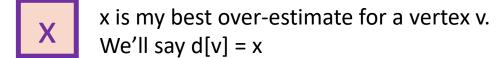


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



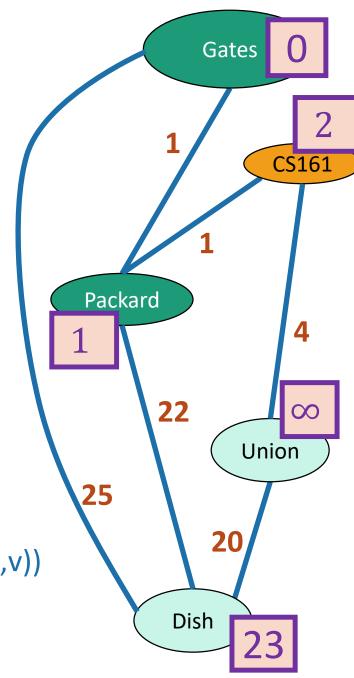


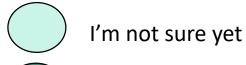


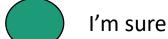


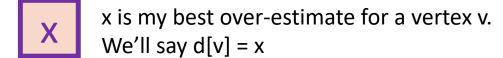


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



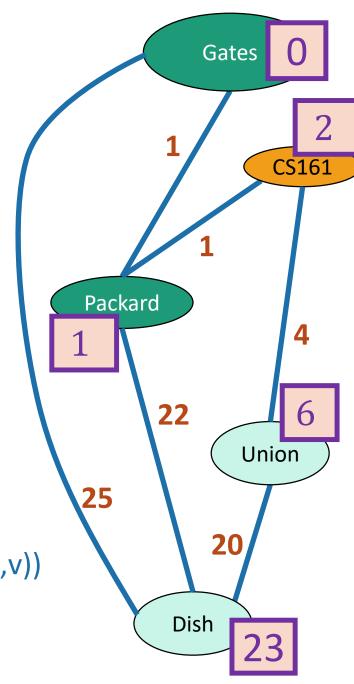


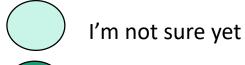


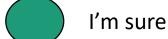


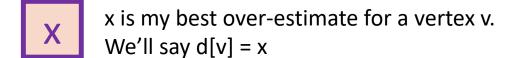


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



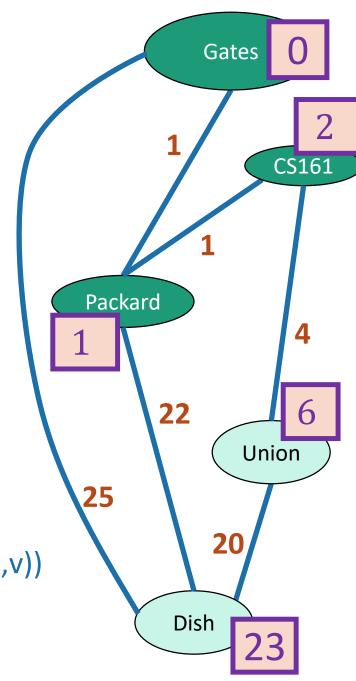


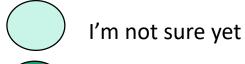


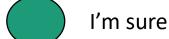


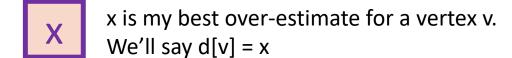


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



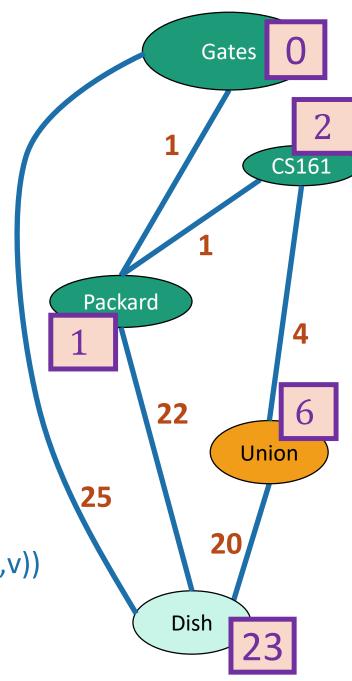


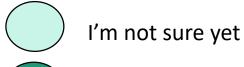


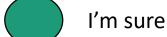


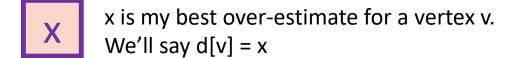


- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat



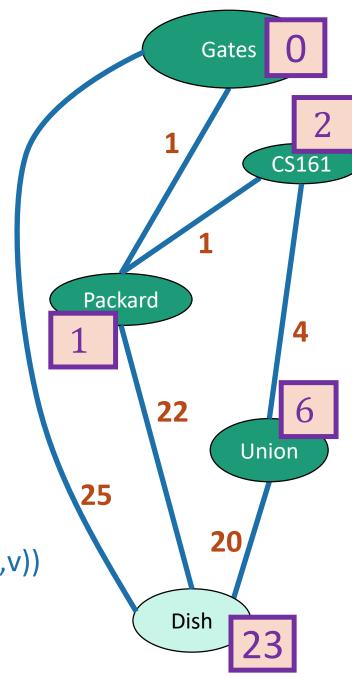


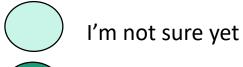








- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat














- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat

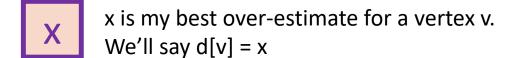


Dijkstra by example

How far is a node from Gates?



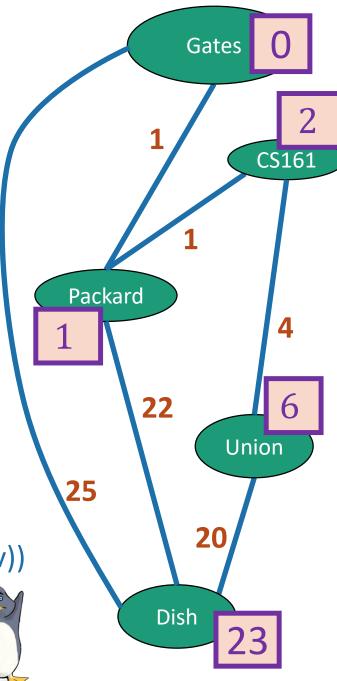






- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] = min(d[v], d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat

More formal pseudocode on board (or see CLRS)!



Why does this work?

Theorem:

- Run Dijkstra on G = (V,E).
- At the end of the algorithm, the estimate d[v] is the actual distance d(Gates,v).
 Let's rename "Gates" to "s", our starting vertex.

Proof outline:

- Claim 1: For all v, $d[v] \ge d(s,v)$.
- Claim 2: When a vertex v is marked sure, d[v] = d(s,v).

prove these

- Claims 1 and 2 imply the theorem.
 - d[v] never increases, so Claims 1 and 2 imply that d[v] weakly decreases until d[v] = d(s,v), then never changes again.
 - By the time we are **sure** about v, d[v] = d(s,v). (Claim 1 again)
 - All vertices are eventually sure. (Stopping condition in algorithm)
 - So all vertices end up with d[v] = d(s,v).

 $d[v] \ge d(s,v)$ for all v.

- Inductive hypothesis.
 - After t iterations of Dijkstra, $d[v] \ge d(s,v)$ for all v.
- Base case:
 - At step 0, d(s, s) = 0, and $d(s, v) \le \infty$
- Inductive step: say hypothesis holds for t.
 - Then at step t+1:
 - We pick u; for each neighbor v:
 - $d[v] \leftarrow min(d[v], d[u] + w(u,v)) \ge d(s,v)$


(Details on board)

By induction, $d(s, v) \le d[v]$

$$d(s, v) \le d(s, u) + d(u, v)$$

 $\le d[u] + w(u, v)$
using induction again for d[u]

So the inductive hypothesis holds for t+1, and Claim 1 follows.



When a vertex u is marked sure, d[u] = d(s,u)

- To begin with:
 - The first vertex marked **sure** has d[s] = d(s,s) = 0.
- For t > 0:
 - Suppose that we are about to add u to the sure list.
 - That is, we picked u in the first line here:
 - Pick the **not-sure** node u with the smallest estimate **d[u]**.
 - Update all u's neighbors v:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))
 - Mark u as Sure.
 - Repeat

Temporary definition:

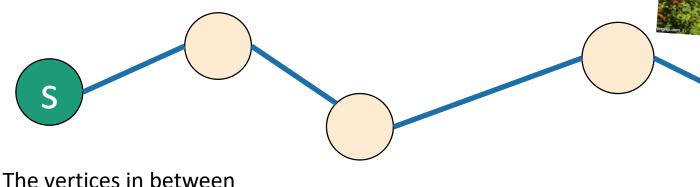
v is "good" means that d[v] = d(s,v)

THOUGHTEXPERIMENT

Claim 2

• Want to show that u is good.


Consider a **true** shortest path from s to u:



are beige because they may or may not be sure.

Temporary definition:

v is "good" means that d[v] = d(s,v)






means not good

"by way of contradiction"

- Want to show that u is good. BWOC, suppose it's not.
- Say z is the last good vertex before u.
- z' is the vertex after z.



Temporary definition:

v is "good" means that d[v] = d(s,v)



means good



means not good

Want to show that u is good. BWOC, suppose it's not.

$$d[z] = d(s, z) \le d(s, u) \le d[u]$$

z is good

This is the shortest path from s to u.

Claim 1

- If d[z] = d[u], then u is good.
- If d[z] < d[u], then z is **sure**.

We chose u so that d[u] was smallest of the unsure vertices.

So assume that z is sure.

z'

It may be that z = s.

It may be that z' = u.

Temporary definition:

v is "good" means that d[v] = d(s,v)

means good

means not good

- Want to show that u is good. BWOC, suppose it's not.
- If z is **sure** then we've already updated z':

•
$$d[z'] \leftarrow \min\{d[z'], d[z] + w(z, z')\}$$
, so

And Z' is 8000d. $d[z'] \leq d[z] + w(z,z') = d(s,z') \leq d[z']$

sub-paths of shortest paths are shortest

CONTRADICTIO

It may be that z = s.

defofupdate

W(Z,Z'

It may be that z' = u.



Temporary definition:

v is "good" means that d[v] = d(s,v)



means good



means not good

Want to show that u is good. BWOC, suppose it's not.

$$d[z] = d(s, z) \le d(s, u) \le d[u]$$

Def. of z

This is the shortest path from s to x

Claim 1

So u is

• If d[z] = d[u], then u is good.

• If d[z] < d[u], then z is **sure**.

aka d[u] = d(s,v)

z

It may be that z = s.

It may be that z' = u.

Back to this slide

Claim 2

When a vertex is marked sure, d[u] = d(s,u)

- To begin with:
 - The first vertex marked sure has d[s] = d(s,s) = 0.
- For t > 0:
 - Suppose that we are about to add u to the sure list.
 - That is, we picked u in the first line here:

Then u is good!

aka d[u] = d(s,u)

- Pick the not-sure node u with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat

Why does this work?

Now back to this slide

• Theorem: At the end of the algorithm, the estimate d[v] is the actual distance d(s,v).

- Proof outline:
 - Claim 1: For all v, $d[v] \ge d(s,v)$.
 - Claim 2: When a vertex is marked sure, d[v] = d(s,v).
- Claims 1 and 2 imply the theorem.
 - We will never mess up d[v] after v is marked sure, because d[v] is a decreasing over-estimate.

Why does this work?

Now back to this slide

Theorem:

- Run Dijkstra on G = (V,E).
- At the end of the algorithm,
 the estimate d[v] is the actual distance d(s,v).

Proof outline:

- Claim 1: For all v, $d[v] \ge d(s,v)$.
- Claim 2: When a vertex v is marked sure, d[v] = d(s,v).


Claims 1 and 2 imply the theorem.

- d[v] never increases, so Claims 1 and 2 imply that d[v] weakly decreases until d[v] = d(s,v), then never changes again.
- By the time we are sure about v, d[v] = d(s,v). (Claim 1 again)
- All vertices are eventually sure. (Stopping condition in algorithm)
- So all vertices end up with d[v] = d(s,v).

What did we just learn?

 Dijkstra's algorithm can find shortest paths in weighted graphs with non-negative edge weights.

- Along the way, it constructs a nice tree.
 - We could post this tree in Gates, and it would be easy for anyone in Gates to figure out what the shortest path is to wherever they want to go.





This is not very precise pseudocode (eg, initialization step is missing)...but it's good enough for this reasoning.

- Pick the **not-sure** node u with the smallest estimate **d[u]**.
- Update all u's neighbors v:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat

This will run for n iterations, since there's one iteration per vertex.

How long does an iteration take?

Depends on how we implement it...

We need a data structure that:

- Stores unsure vertices v
- Keeps track of d[v]
- Can find v with minimum d[v]
 - findMin()
- Can remove that v
 - removeMin(v)
- Can update the d[v]
 - updateKey(v,d)

- Pick the **not-sure** node u
 with the smallest estimate d[u].
- Update all u's neighbors v:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))
- Mark u as Sure.
- Repeat

Total running time is big-oh of:

$$\sum\nolimits_{u \in V} \left(T(findMin) + \left(\sum\limits_{v \in u.neighbors} T(updateKey) \right) + T(removeMin) \right)$$

n(T(findMin) + T(removeMin)) + m T(updateKey)

If we use an array

- T(findMin) = O(n)
- T(removeMin) = O(n)
- T(updateKey) = O(1)

Running time of Dijkstra

```
=O(n(T(findMin) + T(removeMin)) + m T(updateKey))
=O(n^2) + O(m)
=O(n^2)
```

If we use a red-black tree

- T(findMin) = O(log(n))
- T(removeMin) = O(log(n))
- T(updateKey) = O(log(n))

Running time of Dijkstra

```
=O(n(T(findMin) + T(removeMin)) + m T(updateKey))
=O(nlog(n)) + O(mlog(n))
=O((n + m)log(n))
```

Better than an array if the graph is sparse! aka m is much smaller than n²

Is a hash table a good idea here?

Not really:

- Search(v) is fast (in expectation)
- But findMin() will still take time O(n) without more structure.

Can also use a Fibonacci Heap


- This can do all operations in amortized time* O(1).
- Except deleteMin which takes amortized time*
 O(log(n)).
- See CLRS for more!
- This gives (amortized) runtime O(m + nlog(n)) for Dijkstra's algorithm.

*Any sequence of d deleteMin calls takes time at most O(d log(n)). But some of the d may take longer and some may take less time.

Dijkstra is used in practice

- O(nlog(n) + m) is really fast!
- eg, OSPF (Open Shortest Path First), a routing protocol for IP networks, uses Dijkstra.

But there are some things it's not so good at.



Dijkstra Drawbacks

- Needs non-negative edge weights.
- If the weights change, we need to re-run the whole thing.
 - in OSPF, a vertex broadcasts any changes to the network, and then every vertex re-runs Dijkstra's algorithm from scratch.

WE STOPPED HERE IN LECTURE


• Bonus slides follow, but material on the Bellman-Ford algorithm is also in slides for lecture 12.

 The slides below are different than those in lecture 12 (in order to maintain internal consistency within lectures), so they might be interesting for a different perspective.

Bellman-Ford algorithm

Slower than Dijkstra's algorithm

- Can handle negative edge weights.
- Allows for some flexibility if the weights change.
 - We'll see what this means later




Problem

with negative edge weights

- What is the shortest path from Gates to the Union?
- Should still be

- But what about
 - G—P—D—G—P—CS161—Union
- That costs
 - 1-2-3+1+1+4=2.
- And why not



Let's put that aside for a moment



Onwards!

To the Bellman-Ford Bellman-Ford algorithm!

How far is a node from Gates?



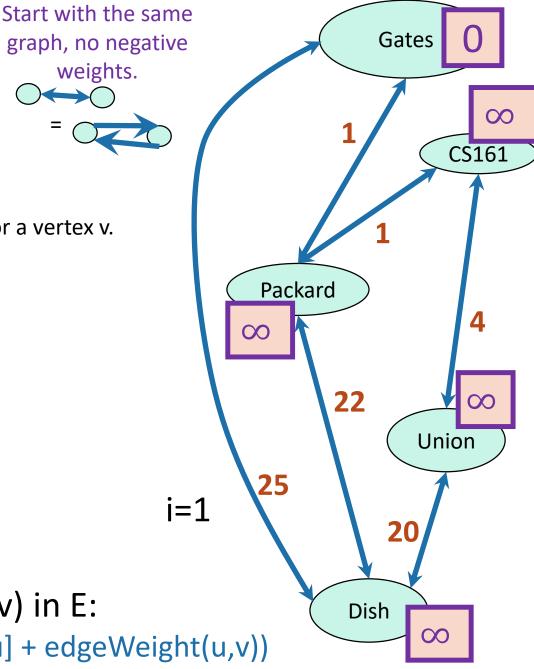
Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

weights.

- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



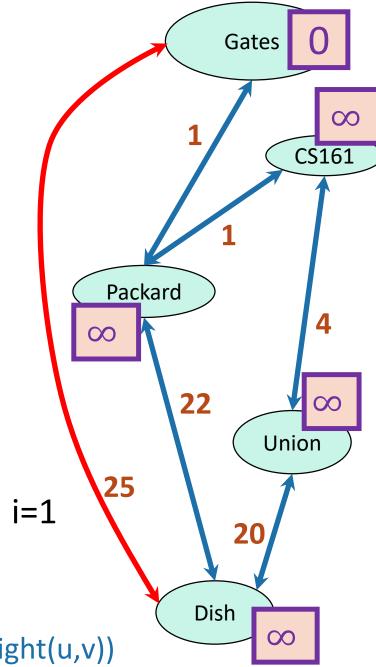
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



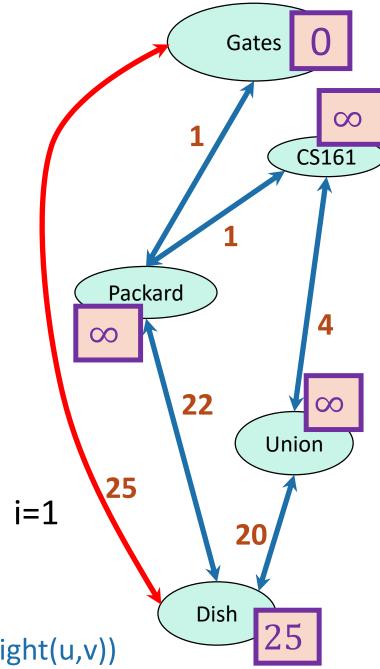
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



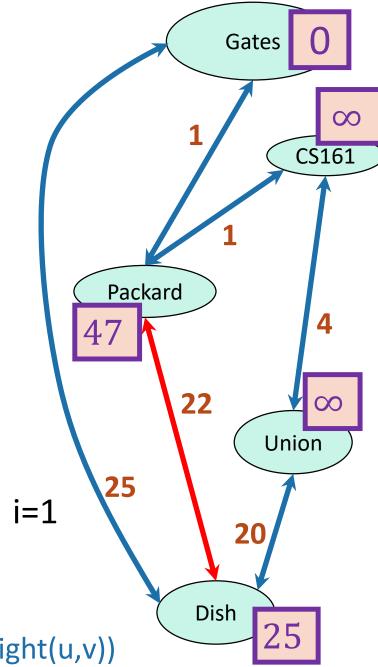
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



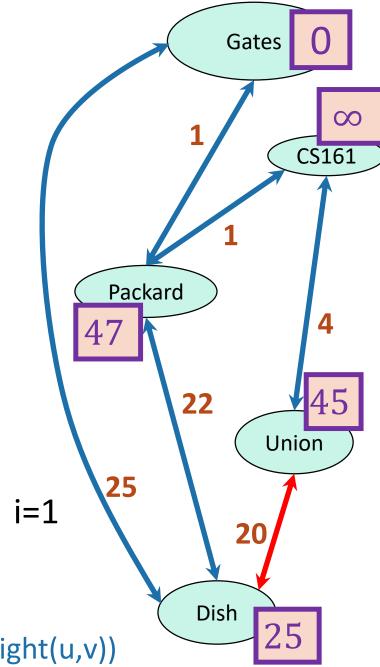
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



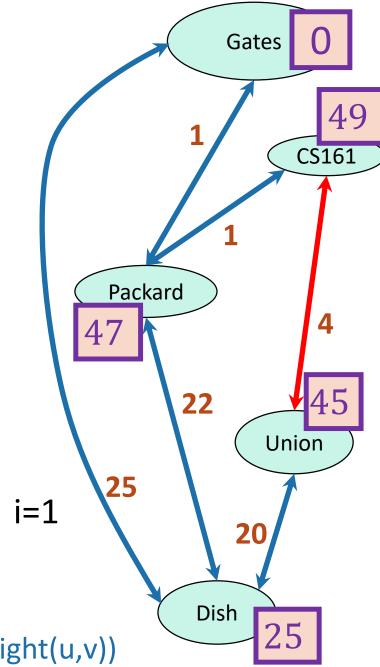
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



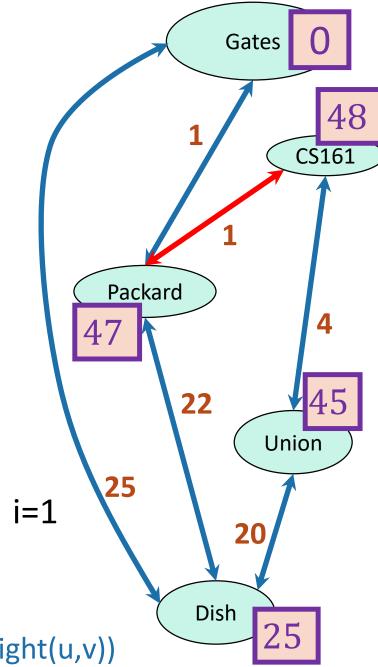
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



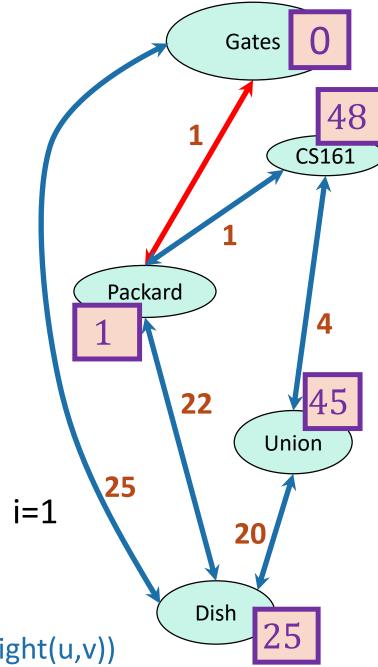
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



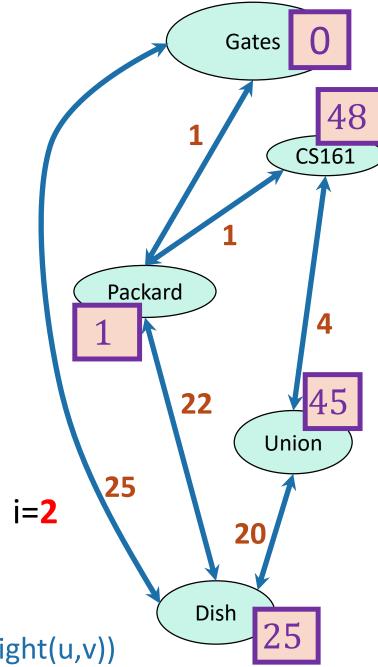
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



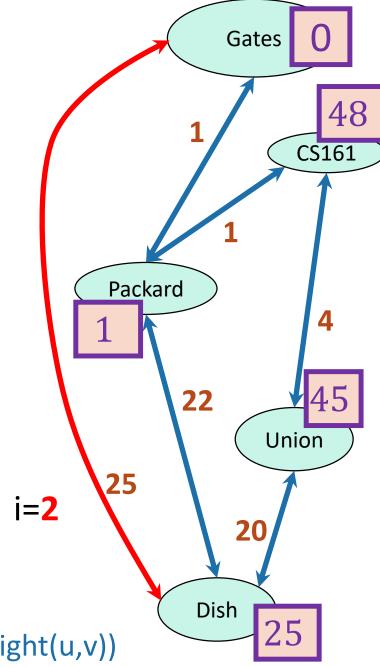
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



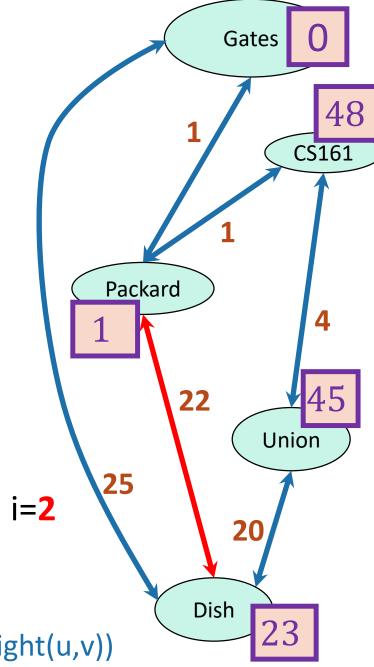
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))



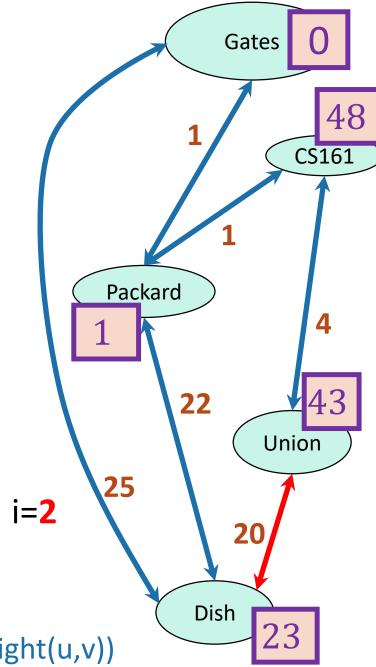
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



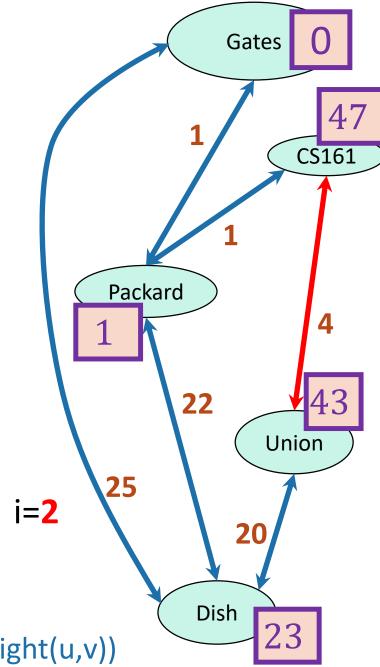
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



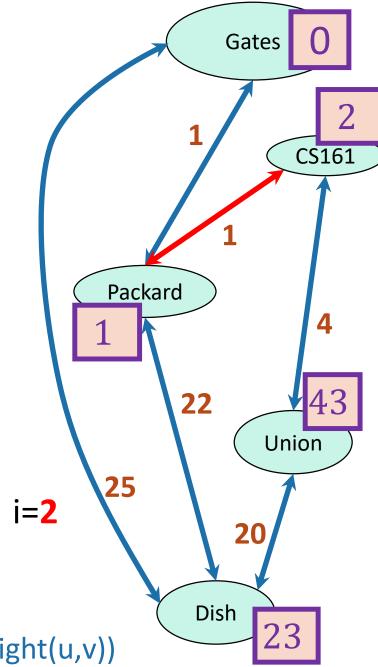
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))



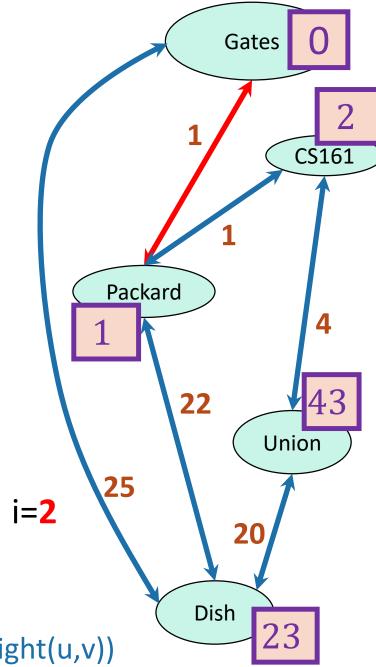
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



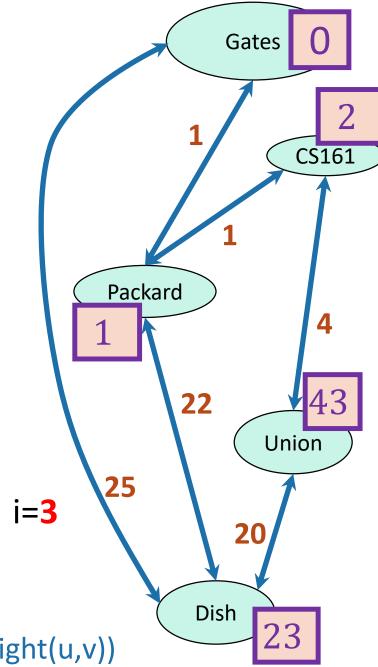
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



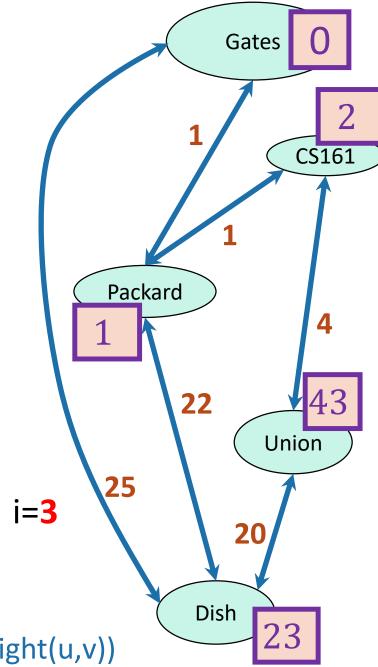
How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))




How far is a node from Gates?



Current edge



- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v] , d[u] + edgeWeight(u,v))



How far is a node from Gates?



Current edge




x is my best over-estimate for a vertex v. We'll say d[v] = x

This will keep on running until i=4, but nothing more will happen.

we say it's converged.

- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- **For** i = 1,...,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))



This seems much slower than Dijkstra

And it is:

Running time O(mn)

- However, it's also more flexible in a few ways.
 - Can handle negative edges
 - If we keep on doing these iterations, then changes in the network will propagate through.

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))

But first

Why does it work as is?

We will show:

- After iteration i, for each v,
 - d[v] is equal to the shortest path between s and v...
 - ...with at most i edges.

In particular:

- After iteration n-1, for each v,
- This is what we want. d[v] is equal to the shortest path between s and v ...
 - ...with at most n-1 euges.

All paths in a graph with n vertices have at most n-1 edges.

Proof by induction

- Inductive Hypothesis:
 - After iteration i, for each v, d[v] is equal to the cost of the shortest path between s and v with at most i edges.
- Base case:
 - After iteration 0... 🔰



Inductive step:

Inductive step

Hypothesis: After iteration i, for each v, d[v] is equal to the cost of the shortest path between s and v with at most i edges.

- Suppose the inductive hypothesis holds for i.
- We want to establish it for i+1.

 Say this is the shortest path between s and v of with at most i+1 edges:

 Let u be the vertex right before v in this path.


 U

 W(u, v)

• By induction, d[u] is the cost of a shortest path between s and u of i edges.

at most i edges

- By setup, d[u] + w(u,v) is the cost of a shortest path between s and v of i+1 edges.
- In the i+1'st iteration, when (u,v) is active, we ensure d[v] <= d[u] + w(u,v).
- So d[v] <= cost of shortest path between s and v with i+1 edges.
- But d[v] = cost of a particular path of at most i+1 edges >= cost of shortest path.
- So d[v] = cost of shortest path with at most i+1 edges.



Proof by induction

- Inductive Hypothesis:
 - After iteration i, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most i edges.
- Base case:
 - After iteration 0...
- Inductive step:
- Conclusion:
 - After iteration n-1, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most n-1 edges.
 - Aka, d[v] = d(s,v) for all v.

Something is wrong

 We never used that there weren't any negative cycles!!



Proof by induction



Inductive Hypothesis:

• After iteration i, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most i edges.

Base case:

• After iteration 0...



Inductive step:

Conclusion:



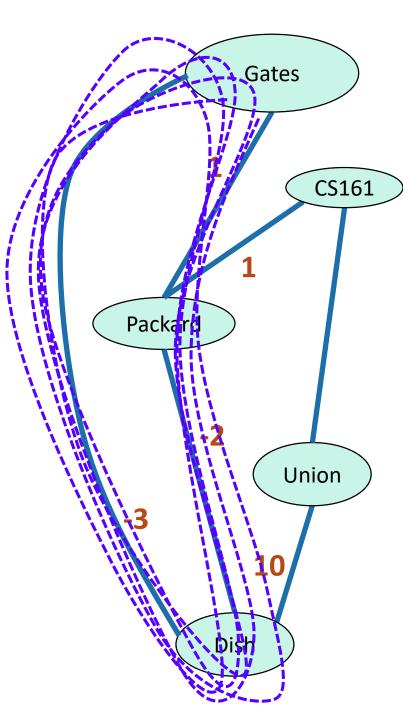
• After iteration n-1, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most n-1 edges.

• Aka, d[v] = d(s,v) for all v



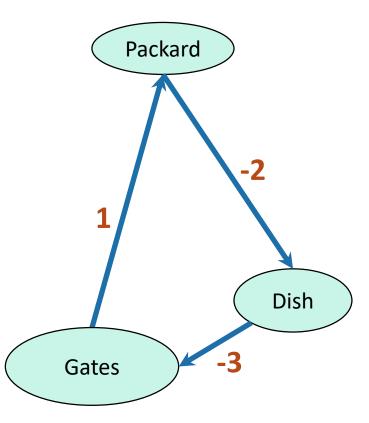
Some paths have more than n-1 edges.

- So we've correctly concluded:
 - After iteration n-1, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most n-1 edges.
- But that's not what we wanted to show.



This is a problem if there are negative cycles.

- A negative cycle is a cycle so that the sum of the edges is negative:
- If there is a negative cycle in G, then there are always shorter paths of length >n
 - Because we can always make a path shorter by going around the cycle.
- We kind of want to ignore this case, though, because "shortest path" doesn't even make sense...



Suppose there are no negative cycles.

- Then all shortest paths are simple paths.
 - A simple path has no cycles.
- It's true that all **simple** paths on n vertices have length at most n-1.

So then we can make the conclusion that we want.

Proof by induction

Suppose there are no negative cycles

• Inductive Hypothesis:


• After iteration i, for each v, d[v] is equal to the cost of the shortest path between s and v of length at most i edges.

Base case:

• After iteration 0...



• Conclusion:



• Aka, the d[v] = d(s,v).



Theorem

- The Bellman-Ford algorithm runs in time O(nm) on a graph G with n vertices and m edges.
- If there are no negative cycles in G, then the BF algorithm terminates with d[v] = d(s,v).


Notice, negative weights are okay.

Okay, so what if there are negative cycles?





- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))





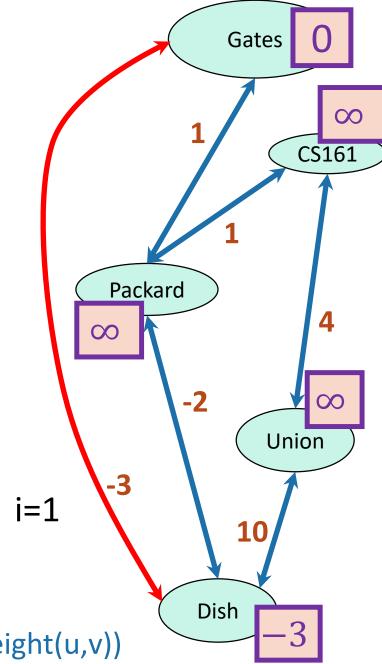


x is my best over-estimate for a vertex v. We'll say d[v] = x

• **For** i = 1,..,n-1:

• For each edge e = (u,v) in E:

d[v] ← min(d[v], d[u] + edgeWeight(u,v))





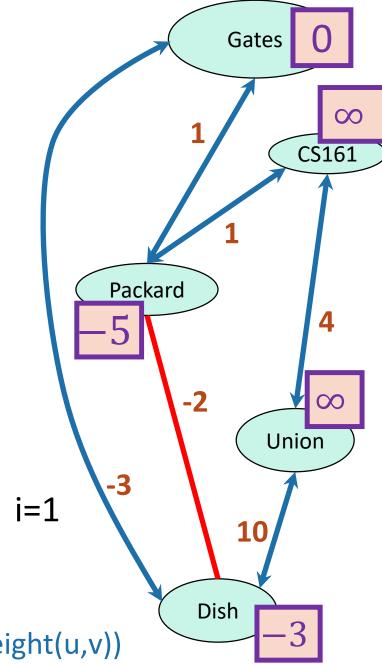


x is my best over-estimate for a vertex v. We'll say d[v] = x

• **For** i = 1,..,n-1:

• For each edge e = (u,v) in E:

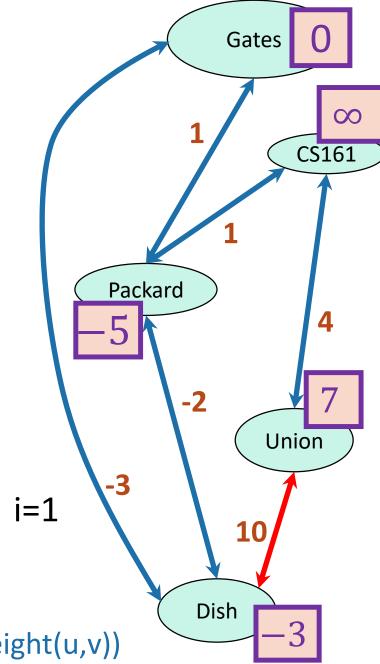
d[v] ← min(d[v], d[u] + edgeWeight(u,v))

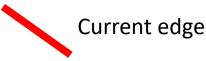






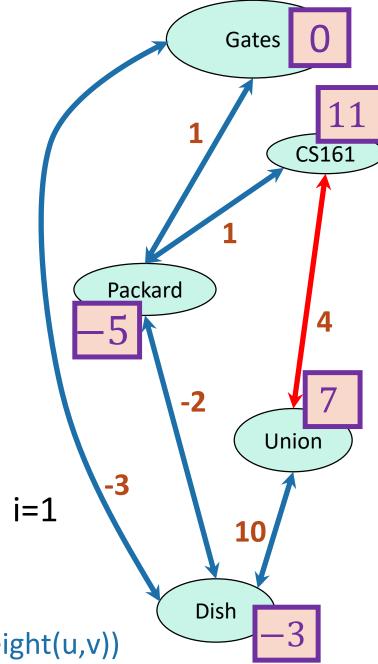
- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))








- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))








- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))







- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Х

x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))






Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))





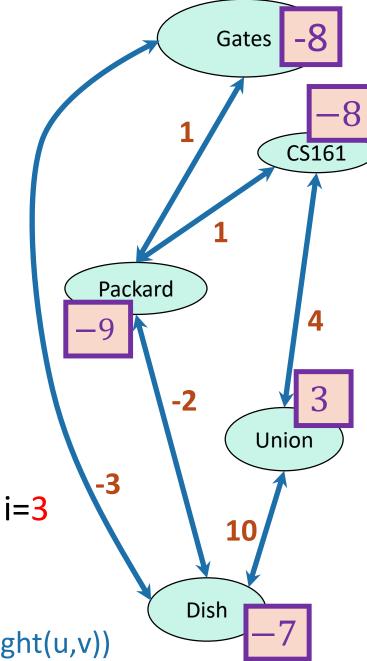
Current edge



x is my best over-estimate for a vertex v. We'll say d[v] = x

You can see where this is going: this will never converge.

- **For** i = 1,..,n-1:
 - **For** each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))





Current edge

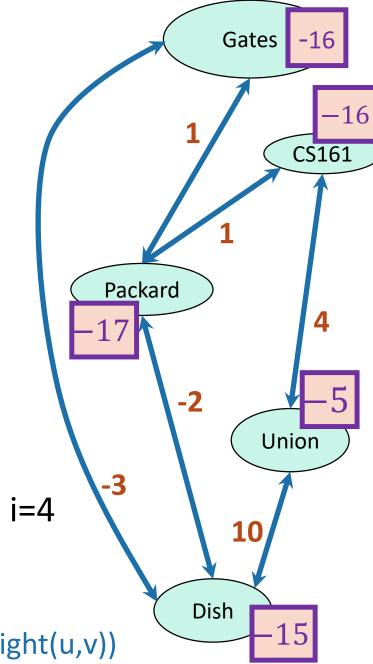


x is my best over-estimate for a vertex v. We'll say d[v] = x

You can see where this is going: this will never converge.

After n-1 iterations, we stop and get something like this.

- **For** i = 1,..,n-1:
 - For each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + edgeWeight(u,v))



How can we tell that this didn't work?

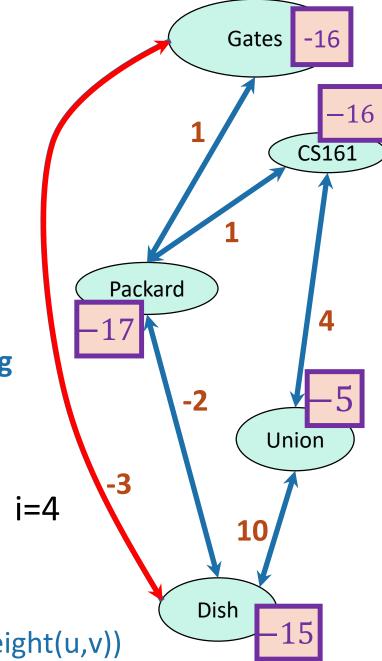
 If we had converged and the algorithm had worked, if we kept going to i=n, nothing would happen

 But if we keep going, then something does happen.

• **For** i = 1,..,n-1:

• For each edge e = (u,v) in E:

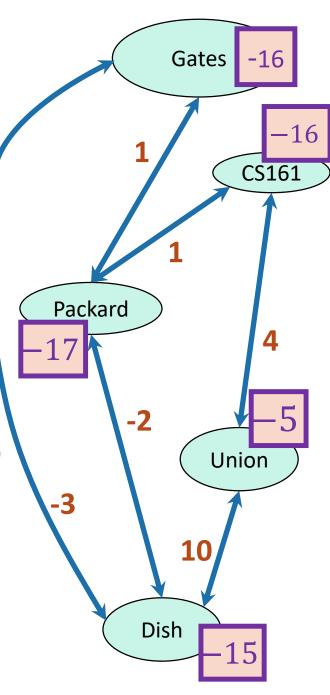
d[v] ← min(d[v], d[u] + edgeWeight(u,v))



This suggests:

Bellman-Ford Algorithm:

- For v in V:
 - $d[v] = \infty$
- d[s] = 0
- For i = 1,...,n-1:
 - **For** each edge e = (u,v) in E:
 - d[v] ← min(d[v], d[u] + weight(u,v))
- For each edge e = (u,v) in E:
 - **if** d[v] < d[u] + weight(u,v)):
 - return negative cycle



What have we just learned?

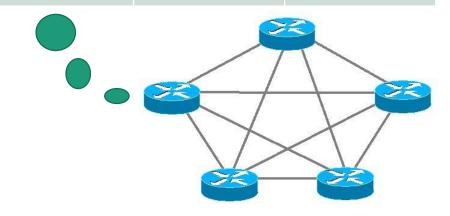
Theorem

- The Bellman-Ford algorithm runs in time O(nm) on a graph G with n vertices and m edges.
- If there are no negative cycles in G, then the BF algorithm terminates with d[v] = d(s,v).
- If there are negative cycles in G, then the BF algorithm returns negative cycle.

Bellman-Ford is also used in practice.

- eg, Routing Information Protocol (RIP) uses something like Bellman-Ford.
 - Older protocol, not used as much anymore.
- Each router keeps a table of distances to every other router.
- Periodically we do a Bellman-Ford update.
- This means that if there are changes in the network, this will propagate. (maybe slowly...)

Destination	Cost to get there	Send to whom?
172.16.1.0	34	172.16.1.1
10.20.40.1	10	192.168.1.2
10.155.120.1	9	10.13.50.0



Recap: shortest paths

- BFS can do it in unweighted graphs
- In weighted graphs:
 - Dijkstra's algorithm is real fast but:
 - doesn't work with negative edge weights
 - is very "centralized"
 - The Bellman-Ford algorithm is slower but:
 - works with negative edge weights
 - can be done in a distributed fashion, every vertex using only information from its neighbors.

Mini-topic (if time) Amortized analysis!

 We mentioned this when we talked about implementing Dijkstra.

*Any sequence of d deleteMin calls takes time at most O(d log(n)). But some of the d may take longer and some may take less time.

 What's the difference between this notion and expected runtime?

Example

Incrementing a binary counter n times.

- Say that flipping a bit is costly.
 - Above, we've noted the cost in terms of bit-flips.

Example

Incrementing a binary counter n times.

- Say that flipping a bit is costly.
 - Some steps are very expensive.
 - Many are very cheap.
- Amortized over all the inputs, it turns out to be pretty cheap.
 - O(n) for all n increments.

This is different from expected runtime.

• The statement is deterministic, no randomness here.



- But it is still weaker than worst-case runtime.
 - We may need to wait for a while to start making it worth it.

Recap

- BFS can do it in unweighted graphs
- In weighted graphs:
 - Dijkstra's algorithm
 - The Bellman-Ford algorithm
- One can implement Dijkstra's algorithm using a fancy data structure (a Fibonacci heap) so that it has good amortized time, O(m + nlog(n)).
 - And now we have a slightly better idea what amortized time means.